

The orthodox framework of equity valuation rests on the premise that the intrinsic value of a share represents the present value of all future cash flows attributable to that share, discounted at a rate that reflects both the time value of money and a compensation for bearing systematic risk. This framework — most formally articulated in the Dividend Discount Model and its generalisations — assumes that rational agents process all available information efficiently, thereby precluding sustained arbitrage opportunities. Yet decades of empirical evidence have mounted a formidable challenge to this elegant theoretical construct, exposing persistent anomalies that a strictly neoclassical model struggles to reconcile.

The equity risk premium — defined as the excess return demanded by investors for holding equities over a nominally risk-free instrument such as a short-dated government security — has historically averaged between four and six percent annually in developed markets. This magnitude has long puzzled theorists, constituting what Mehra and Prescott famously termed the "equity premium puzzle": under plausible parameterisations of investor risk aversion, standard consumption-based asset pricing models generate equity premia far smaller than those observed empirically. Proposed resolutions have ranged from habit formation in preferences and disaster risk incorporation, to the liquidity premium hypothesis, which posits that the observed premium partly compensates for the relative illiquidity and transaction cost burden of equity relative to sovereign debt.

Compounding the valuation challenge is the cyclically adjusted price-to-earnings ratio, popularised by economist Robert Shiller, which smooths reported corporate earnings across a ten-year window to attenuate the distorting effects of business cycle fluctuations on single-period earnings figures. The CAPE ratio has demonstrated meaningful predictive power over decade-long return horizons: periods characterised by elevated CAPE valuations have systematically preceded lower prospective real returns, while troughs in CAPE have corresponded with subsequent above-average performance. Critics, however, contend that structural changes in accounting standards, sectoral composition shifts — particularly the ascendancy of asset-light technology companies whose earnings quality and capital intensity differ fundamentally from prior industrial cohorts — and the secular decline in real interest rates since the 1980s collectively impair the ratio's cross-temporal comparability.

The efficient market hypothesis, in its semi-strong form, asserts that current market prices fully reflect all publicly available information, rendering consistent alpha generation through fundamental analysis theoretically impossible. Paradoxically, this hypothesis depends on the very existence of active arbitrageurs willing to incur research costs to identify and exploit mispricings — a contradiction elegantly captured in the Grossman-Stiglitz paradox. If markets were perfectly efficient, no rational agent would expend resources on price discovery, thereby

causing efficiency itself to erode. The equilibrium condition, therefore, implies markets that are approximately but not perfectly efficient, leaving a residual layer of exploitable inefficiency that varies inversely with the costs and risks of arbitrage.

Momentum, mean reversion, and the size and value factors documented in the Fama-French three-factor model further complicate the single-factor Capital Asset Pricing Model's explanatory framework. Momentum — the empirical regularity whereby recent outperformers tend to continue outperforming over intermediate horizons of three to twelve months — is particularly difficult to reconcile with rational expectations, and has been attributed variously to investor underreaction to earnings information, herding behaviour among institutional allocators, and feedback trading dynamics. Mean reversion, operating on longer horizons of three to five years, ostensibly reflects the correction of prior overvaluation and is consistent with contrarian investment strategies that systematically purchase depressed assets and divest appreciated ones.

Liquidity risk, while often conflated with market risk in practitioner discourse, represents a conceptually distinct dimension of investment risk that commands a separate premium in equilibrium. Assets with thin secondary markets — characterised by wide bid-ask spreads, high price impact per unit traded, and intermittent trading activity — impose holding costs on investors that extend beyond volatility-based risk measures. During episodes of systemic stress, liquidity premia exhibit sharp non-linearities: the flight-to-quality phenomenon drives capital simultaneously toward highly liquid sovereign instruments and away from illiquid credit, small-capitalisation equities, and alternative assets, generating correlated drawdowns precisely when diversification is most needed.

Behavioural finance has further eroded the epistemic confidence of the rational agent model by cataloguing systematic cognitive biases that distort investor decision-making in reproducible ways. Loss aversion — the asymmetric weighting of potential losses relative to equivalent gains, formalised in Kahneman and Tversky's prospect theory — induces suboptimal holding periods, as investors prematurely realise gains to lock in the psychological satisfaction of being correct, while holding losing positions beyond rational thresholds to avoid the discomfort of crystallising a loss. Overconfidence biases, confirmation heuristics, and the availability cascade generated by salient recent market events further compound the gap between theoretical and observed market behaviour.

In synthesis, the valuation of equities in contemporary markets defies reduction to any single explanatory paradigm. The interplay of rational discounting, risk premia for systematic and idiosyncratic exposures, liquidity constraints, and pervasive cognitive distortions generates a dynamic and non-stationary market environment in which no static model can claim universal validity. Practitioners who acknowledge this irreducible complexity — maintaining epistemic

humility about the precision of their valuation estimates, constructing portfolios robust to multiple plausible scenarios, and accounting explicitly for liquidity and behavioural tail risks — are, on the evidence of both academic literature and practitioner experience, best positioned to navigate the enduring uncertainty that defines equity markets.